




राष्ट्रीय प्रौद्योगिकी संस्थान सिक्किम
NATIONAL INSTITUTE OF TECHNOLOGY SIKKIM

Deep Learning: Feed Forward Neural Network

Course Instructor:
Dr. Bam Bahadur Sinha
Assistant Professor
Computer Science & Engineering
National Institute of Technology
Sikkim

What is the
**LOSS
FUNCTION**
that you use for
a multi-class
classification
problem?



$$b = [0 \ 0]$$

$$W_1 = \begin{bmatrix} 0.1 & 0.3 & 0.8 & -0.4 \\ -0.3 & -0.2 & 0.5 & 0.5 \\ -0.3 & 0.1 & 0.5 & 0.4 \\ 0.2 & 0.5 & -0.9 & 0.7 \end{bmatrix}$$

$$W_2 = \begin{bmatrix} 0.3 & 0.8 & -0.2 & -0.4 \\ 0.5 & -0.2 & -0.3 & 0.5 \\ 0.3 & 0.1 & 0.6 & 0.6 \end{bmatrix}$$

$$x = [0.3 \quad -0.4 \quad 0.6 \quad 0.2] \quad y = [0 \quad 0 \quad 1]$$

Output :

$$a_1 = W_1 * x + b_1 = [0.31 \quad 0.39 \quad 0.25 \quad -0.54]$$

$$h_1 = \text{sigmoid}(a_1) = [0.58 \quad 0.60 \quad 0.56 \quad 0.37]$$

$$a_2 = W_2 * h_1 + b_2 = [0.39 \quad 0.18 \quad 0.79]$$


$$\hat{y} = \text{softmax}(a_2) = [0.3024 \quad 0.2462 \quad 0.4514]$$

Cross Entropy Loss:

$$L(\Theta) = - \sum_{i=1}^k y_i \log(\hat{y}_i)$$

$$L(\Theta) = -1 * \log(0.4514) \\ = 0.7954$$

Multi-class classification problem



Given weights, we know how to compute the model's output for a given input

Given weights, we know how to compute the model's loss for a given input

But, who will give us the weights ?

Learning Algorithm

Learning Algorithm

Initialise w, b

Iterate over data:


compute \hat{y}

compute $\mathcal{L}(w, b)$

$w = w - \eta \Delta w$


$b = b - \eta \Delta b$

till satisfied



Earlier : w, b

Now : w_{111}, w_{112}, \dots



Earlier : $L(w, b)$

Now : $L(w_{111}, w_{112}, \dots)$

Gradient Descent Learning Algorithm for Multi-class Classification problem

Initialise ~~w, b~~ w_{111}, w_{112}

Iterate over data:

compute \hat{y}

compute $\mathcal{L}(w, b)$

$$w_{111} = w_{111} - \eta \Delta w_{111}$$

$$w_{112} = w_{112} - \eta \Delta w_{112}$$

....

$$w_{313} = w_{313} - \eta \Delta w_{313}$$

till satisfied

How do you
check the
performance
of a Deep
Neural
Network?

(Binary
Classification)

Indian Liver Patient Records * - whether person needs to be diagnosed or not ?

Test Data

Age	Albumin	T_Bilirubin
65	3.3	0.7
62	3.2	10.9
20	4	1.1
84	3.2	0.7

...

y	Predicted
0	0
0	1
1	1
1	0



$$Accuracy = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$

$$= \frac{2}{4} = 50\%$$

How do you
check the
performance
of a Deep
Neural
Network?

(Multi-class
Classification)

Test Data

0
1
3
5
1

y	Predicted
0	0
1	7
3	8
5	5
1	1

$$Accuracy = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$

$$= \frac{3}{5} = 60\%$$

Summary and Roadmap



Real inputs

$$x_i \in \mathbb{R}$$



Squared Error Loss :

$$L(\Theta) = \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^d (\hat{y}_{ij} - y_{ij})^2$$

Cross Entropy Loss:

$$L(\Theta) = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^d y_{ij} \log(\hat{y}_{ij})$$




- Binary Classification
- Multi-class classification
- Regression



Gradient Descent **with**
backpropagation




$$\hat{y} = \frac{1}{1 + e^{-(w_{21} * (\frac{1}{1 + e^{-(w_{11} * x_1 + w_{12} * x_2 + b_1)}}) + w_{22} * (\frac{1}{1 + e^{-(w_{13} * x_1 + w_{14} * x_2 + b_1)}}) + b_2)}}$$



$$Accuracy = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$

How many derivatives do we need to compute and how do we compute them?



- Let us focus on the highlighted weight (w_{222})
- To learn this weight, we have to compute partial derivative w.r.t loss function

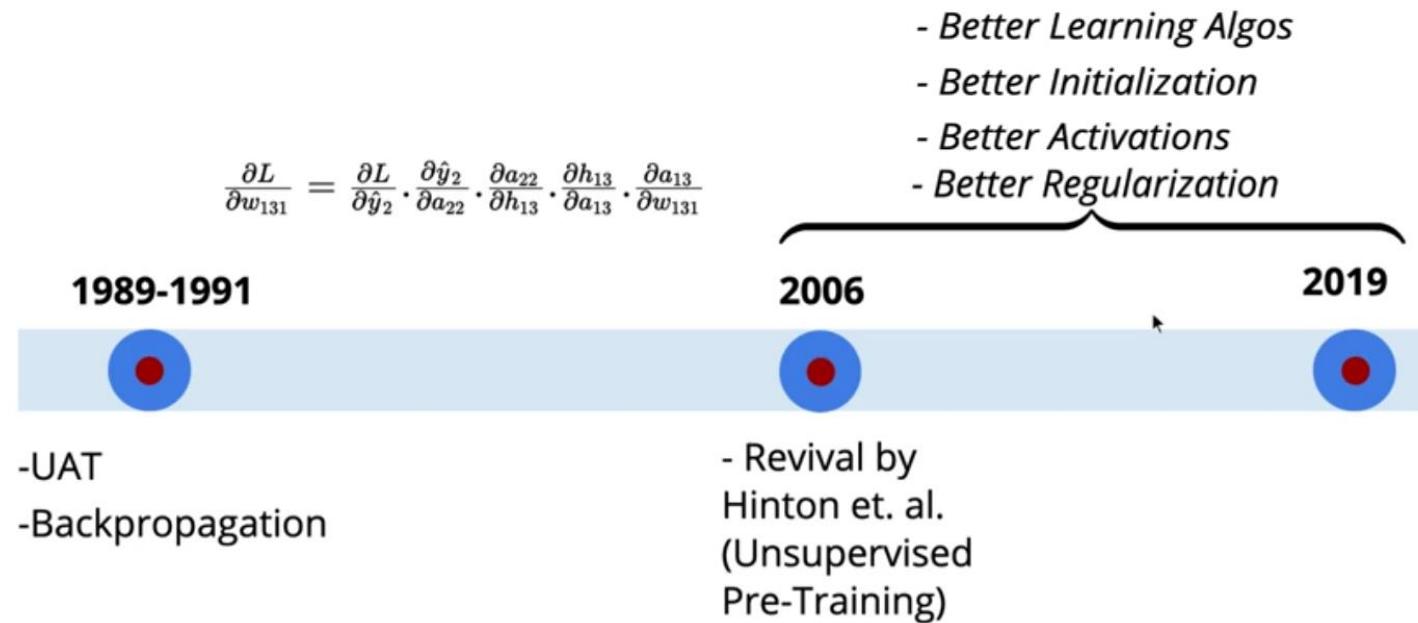
$$(w_{222})_{t+1} = (w_{222})_t - \eta * \left(\frac{\partial L}{\partial w_{222}} \right)$$

$$\begin{aligned} \frac{\partial L}{\partial w_{222}} &= \left(\frac{\partial L}{\partial a_{22}} \right) \cdot \left(\frac{\partial a_{22}}{\partial w_{222}} \right) \\ &= \left(\frac{\partial L}{\partial h_{22}} \right) \cdot \left(\frac{\partial h_{22}}{\partial a_{22}} \right) \cdot \left(\frac{\partial a_{22}}{\partial w_{222}} \right) \\ &= \left(\frac{\partial L}{\partial a_{31}} \right) \cdot \left(\frac{\partial a_{31}}{\partial h_{22}} \right) \cdot \left(\frac{\partial h_{22}}{\partial a_{22}} \right) \cdot \left(\frac{\partial a_{22}}{\partial w_{222}} \right) \\ &= \left(\frac{\partial L}{\partial \hat{y}} \right) \cdot \left(\frac{\partial \hat{y}}{\partial a_{31}} \right) \cdot \left(\frac{\partial a_{31}}{\partial h_{22}} \right) \cdot \left(\frac{\partial h_{22}}{\partial a_{22}} \right) \cdot \left(\frac{\partial a_{22}}{\partial w_{222}} \right) \end{aligned}$$

If DL is not
working
appropriate:

- Better Optimization
- Better Activation Function
- Better Weight Initialization
- Better Regularizer
- Better Compute
- Better Data

Deep Learning Timeline



Better Learning Algorithm

Gradient Descent Update Rule

$$w = w - \eta \frac{\partial \mathcal{L}(w)}{\partial w}$$

- How do we compute the gradients?
- What data should we use for computing the gradients?
- How do we use the gradients?
- Can we come-up with a better update rule?

Gradient Descent Update

```
X = [0.5, 2.5]
Y = [0.2, 0.9]





def f(w, b, x):
    #sigmoid with parameters w, b
    return 1.0 / (1.0 + np.exp(-(w*x + b)))

def error(w, b):
    err = 0.0
    for x, y in zip(X, Y):
        fx = f(w, b, x)
        err += 0.5 * (fx - y) ** 2
    return err

def grad_b(w, b, x, y):
    fx = f(w, b, x)
    return (fx - y) * fx * (1 - fx)

def grad_w(w, b, x, y):
    fx = f(w, b, x)
    return (fx - y) * fx * (1 - fx) * x

def do_gradient_descent():
    w, b, eta = -2, -2, 1.0
    max_epochs = 1000
    for i in range(max_epochs):
        dw, db = 0, 0
        for x, y in zip(X, Y):
            dw += grad_w(w, b, x, y)
            db += grad_b(w, b, x, y)
        w = w - eta * dw
        b = b - eta * db
```





Drawback of Gradient Computation

Initialise w, b randomly.

Iterate over data:


till satisfied



Momentum based Gradient Descent

Issues

It takes a lot of time to navigate regions having gentle slope (because the gradient in these regions is very small)



Intuitive Solution

If I am repeatedly being asked to go in the same direction, then I should probably gain some confidence & start taking bigger steps in that direction.

Mathematical Intuition

Gradient Descent Update Rule

$$w_{t+1} = w_t - \eta \nabla w_t$$

Momentum based Gradient Descent Update Rule

$$v_t = \gamma * v_{t-1} + \eta \nabla w_t$$

$$w_{t+1} = w_t - v_t$$

Momentum Based Gradient Descent

$$v_t = \gamma \cdot v_{t-1} + \eta \nabla w_t$$
$$w_{t+1} = w_t - v_t$$

$$v_0 = 0$$

$$v_1 = \gamma \cdot v_0 + \eta \nabla w_1 = \eta \nabla w_1$$

$$v_2 = \gamma \cdot v_1 + \eta \nabla w_2 = \gamma \cdot \eta \nabla w_1 + \eta \nabla w_2$$

$$v_3 = \gamma \cdot v_2 + \eta \nabla w_3 = \gamma(\gamma \cdot \eta \nabla w_1 + \eta \nabla w_2) + \eta \nabla w_3$$
$$= \gamma \cdot v_2 + \eta \nabla w_3 = \gamma^2 \cdot \eta \nabla w_1 + \gamma \cdot \eta \nabla w_2 + \eta \nabla w_3$$

$$v_4 = \gamma \cdot v_3 + \eta \nabla w_4 = \gamma^3 \cdot \eta \nabla w_1 + \gamma^2 \cdot \eta \nabla w_2 + \gamma \cdot \eta \nabla w_3 + \eta \nabla w_4$$

$$\vdots$$

$$v_t = \gamma \cdot v_{t-1} + \eta \nabla w_t = \gamma^{t-1} \cdot \eta \nabla w_1 + \gamma^{t-2} \cdot \eta \nabla w_2 + \dots + \eta \nabla w_t$$

*Exponential Decaying Average

Gradient Descent Update Rule

Vs

Momentum Based Gradient Descent

```
X = [0.5, 2.5]
Y = [0.2, 0.9]

def f(w, b, x):
    #sigmoid with parameters w, b
    return 1.0 / (1.0 + np.exp(-(w*x + b)))

def error(w, b):
    err = 0.0
    for x, y in zip(X, Y):
        fx = f(w, b, x)
        err += 0.5 * (fx - y) ** 2
    return err

def grad_b(w, b, x, y):
    fx = f(w, b, x)
    return (fx - y) * fx * (1 - fx)

def grad_w(w, b, x, y):
    fx = f(w, b, x)
    return (fx - y) * fx * (1 - fx) * x

def do_gradient_descent():
    w, b, eta = -2, -2, 1.0
    max_epochs = 1000
    for i in range(max_epochs):
        dw, db = 0, 0
        for x, y in zip(X, Y):
            dw += grad_w(w, b, x, y)
            db += grad_b(w, b, x, y)
        w = w - eta * dw
        b = b - eta * db
```

Gradient Descent Update Rule


Momentum^{*}based Gradient Descent Update Rule

$$v_t = \gamma * v_{t-1} + \eta \nabla w_t$$
$$w_{t+1} = w_t - v_t$$


```
def do_momentum_gradient_descent():
    w, b, eta, max_epochs = -2, -2, 1.0, 1000
    v_w, v_b = 0, 0
    for i in range(max_epochs):
        dw, db = 0, 0
        for x, y in zip(X, Y):
            dw += grad_w(w, b, x, y)
            db += grad_b(w, b, x, y)
        v_w = gamma*v_w + eta * dw
        v_b = gamma*v_b + eta * db

        w = w - v_w
        b = b - v_b
```

Momentum Gradient Descent



Momentum Gradient Descent



Gradient Descent

Convergence Graph - GD Vs MGD

Gradient Descent Update Rule