Deep Learning :
Multi-headed Attention in Transformers

राष्ट्रीय प्रौद्योगिकी संस्थान सिक्किम
NATIONAL INSTITUTE OF TECHNOLOGY SIKKIM

Course Instructor:
Dr. Bam Bahadur Sinha
Assistant Professor
Computer Science & Engineering
National Institute of Technology
Sikkim

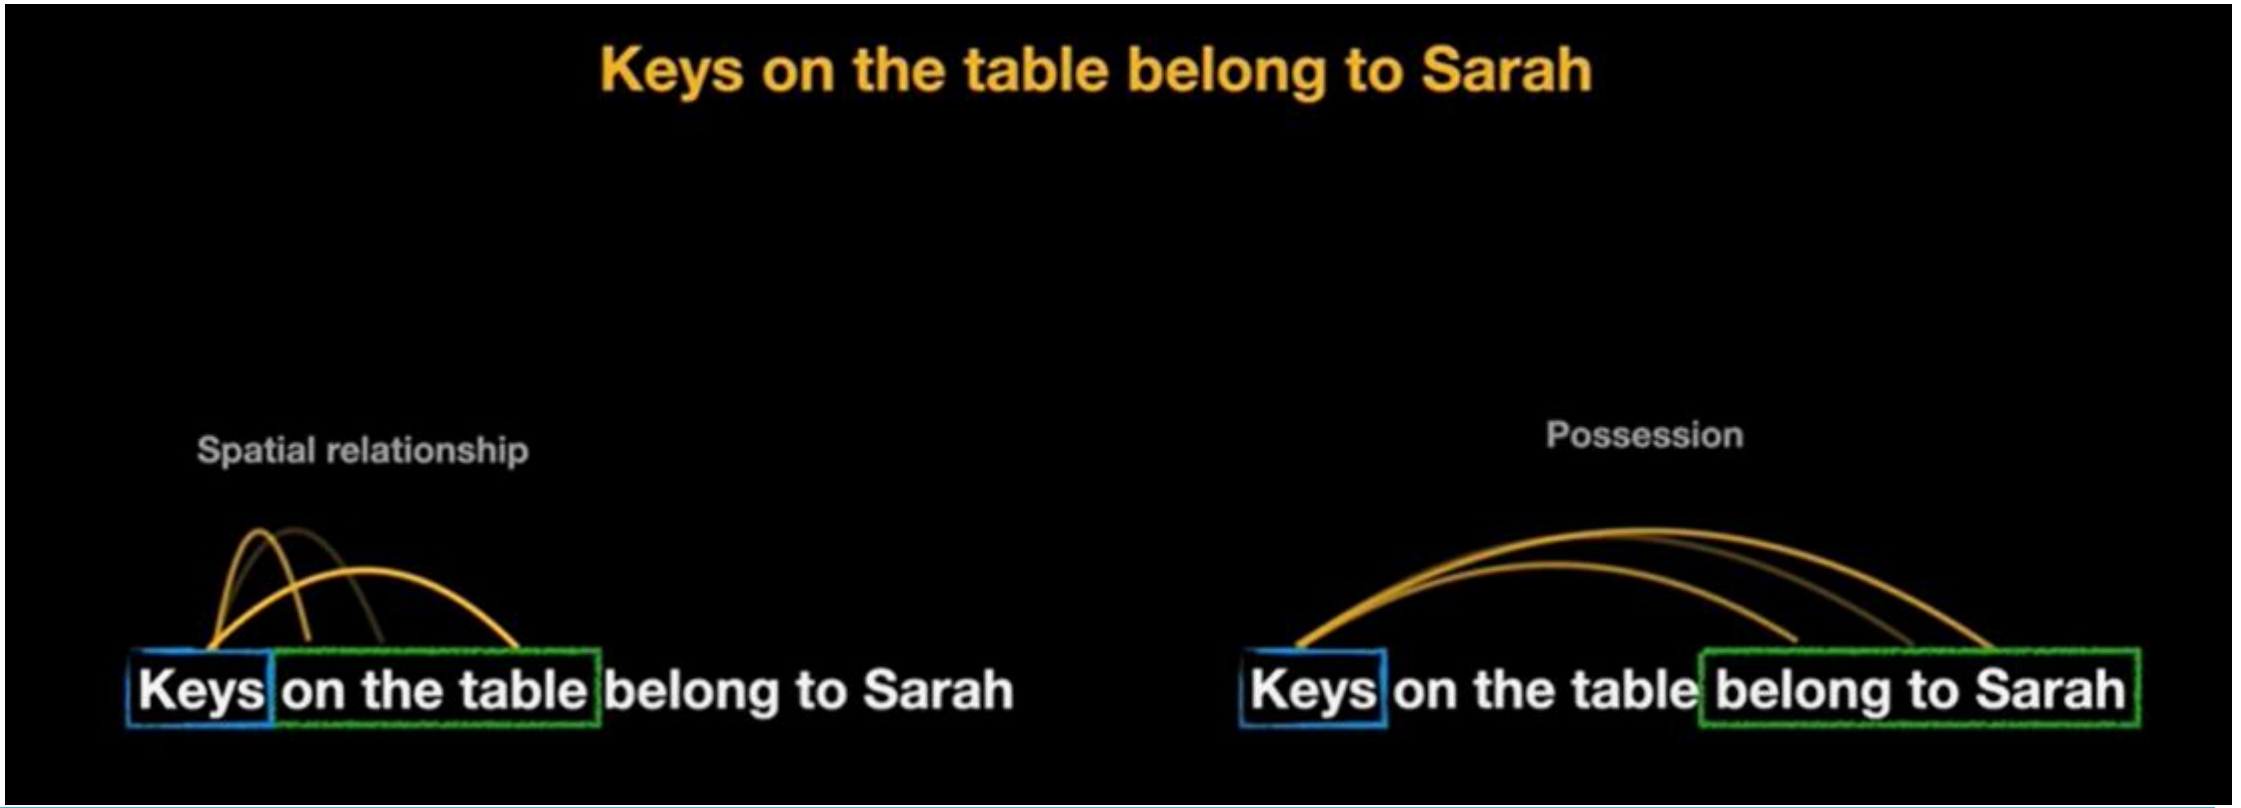

Multi-headed attention

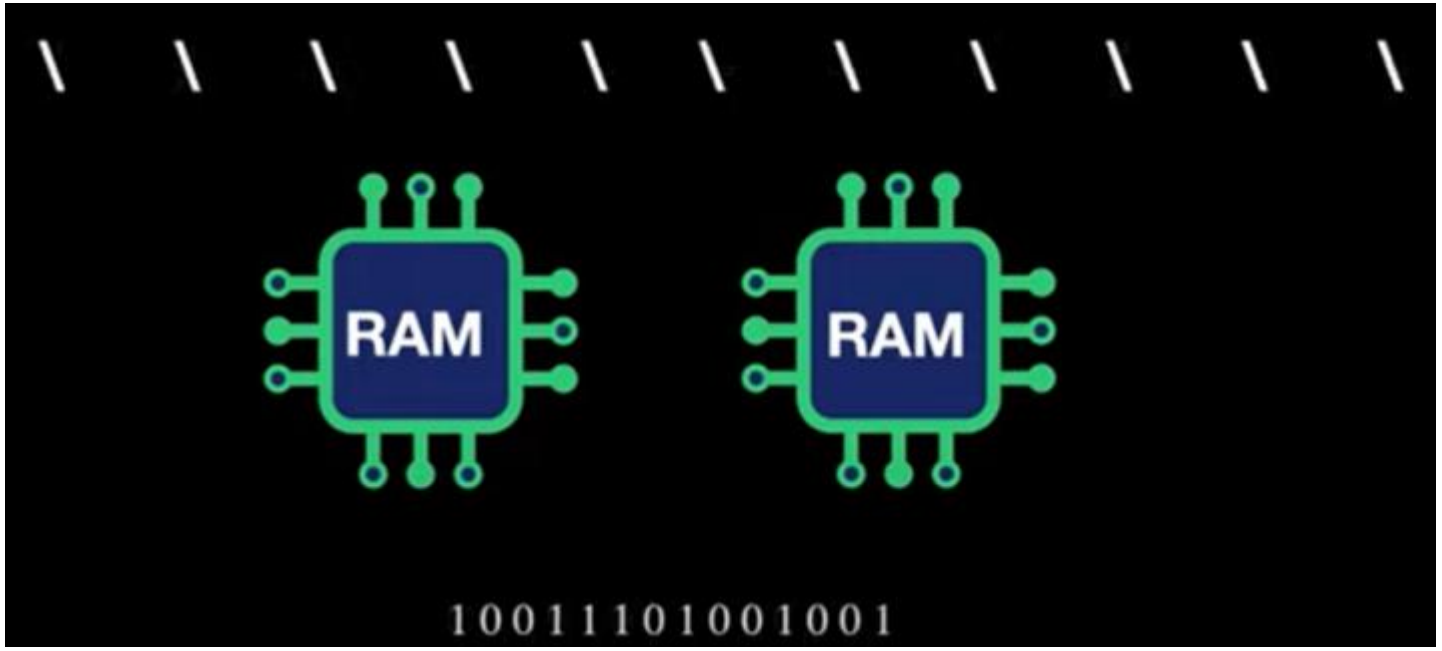

Multi-headed attention

Multi-headed attention

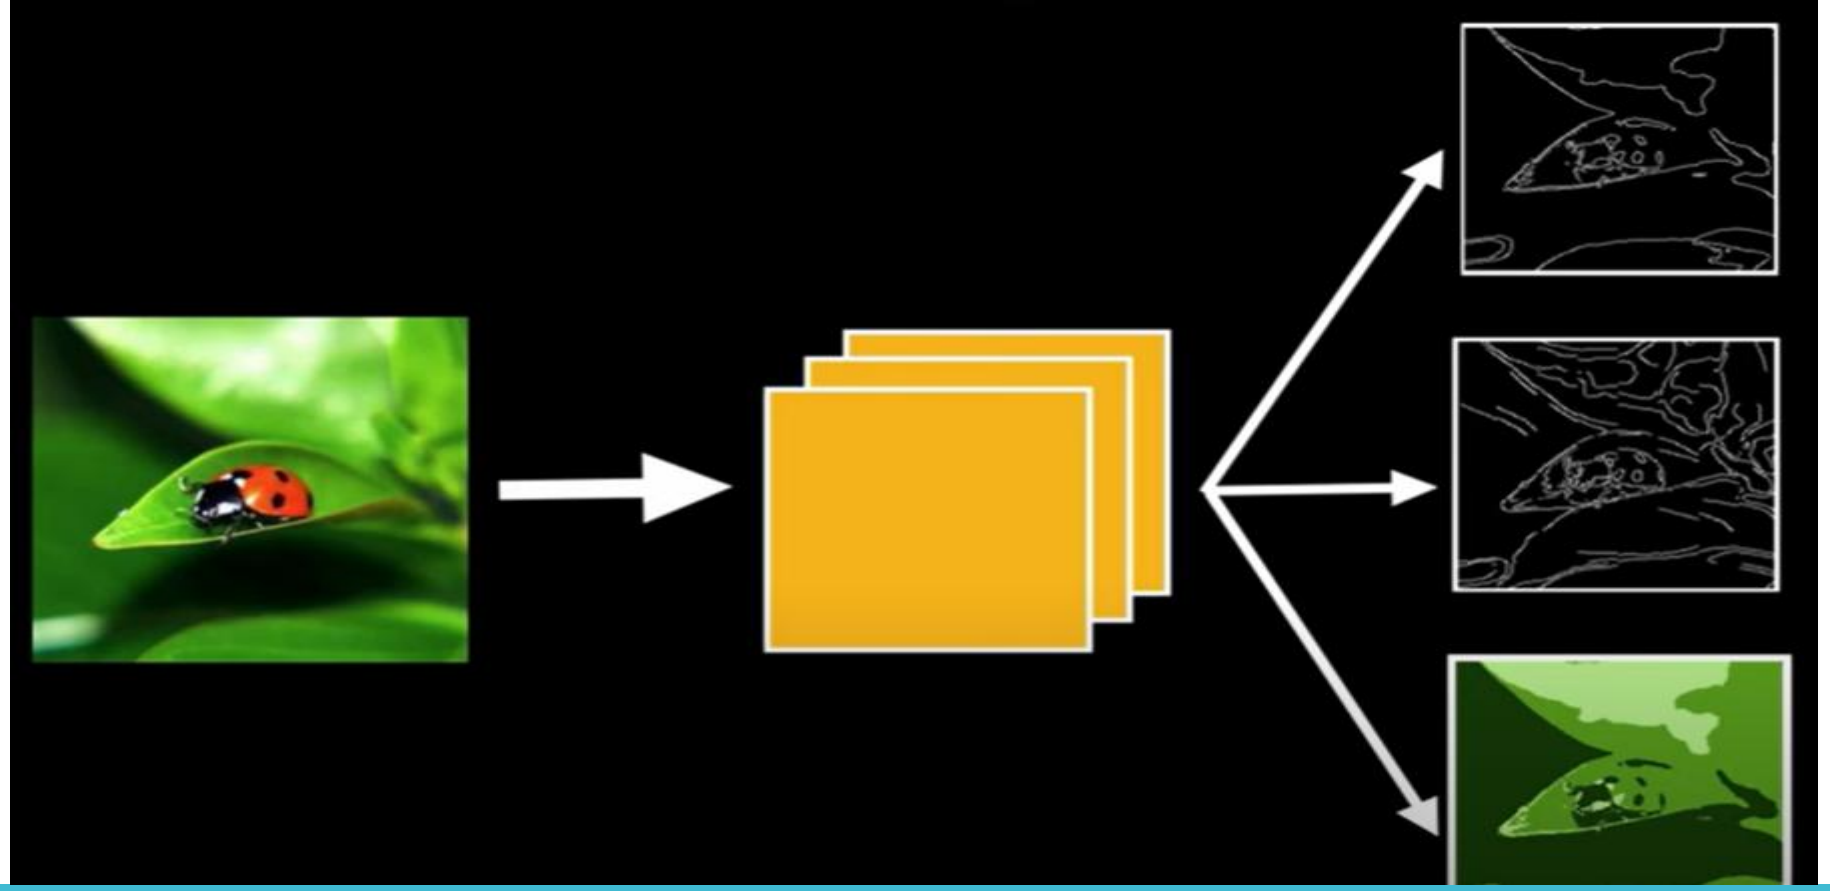

Multi-headed attention (in comparison to CNN)

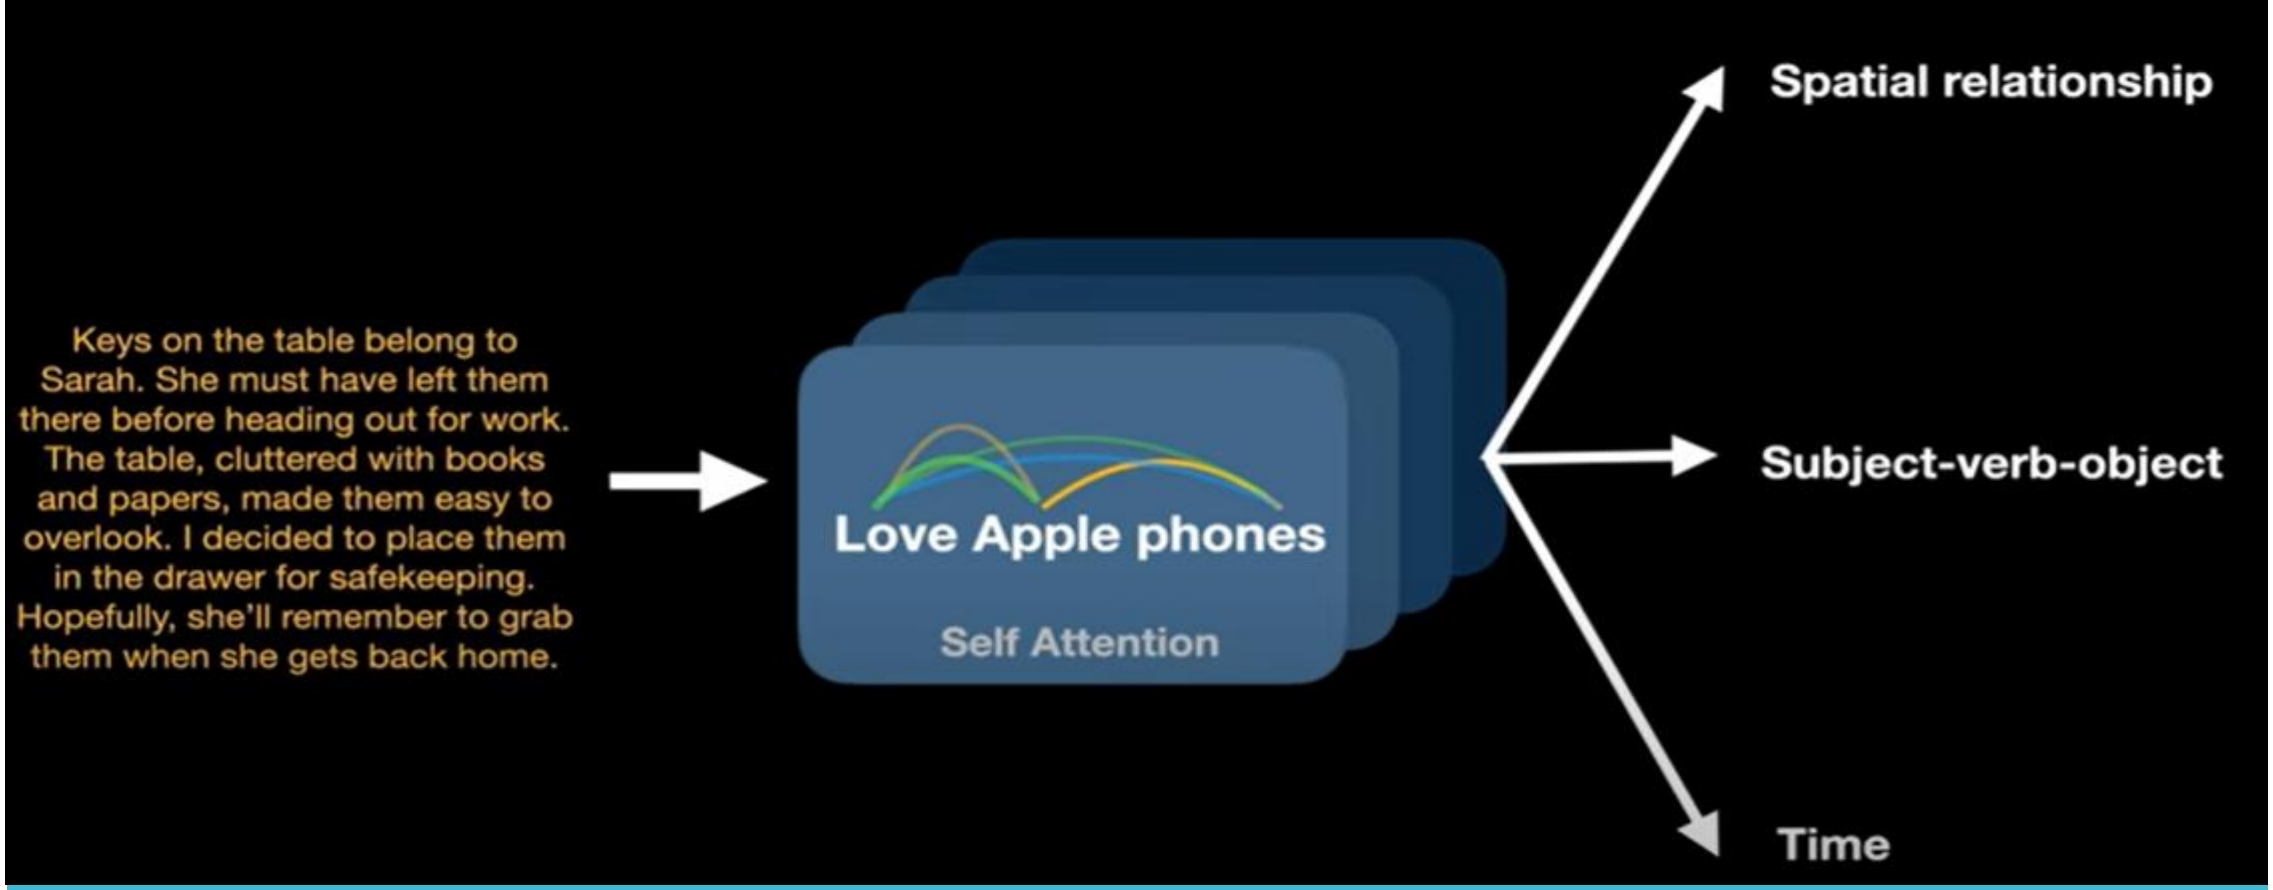

Multi-headed attention (in comparison to CNN)

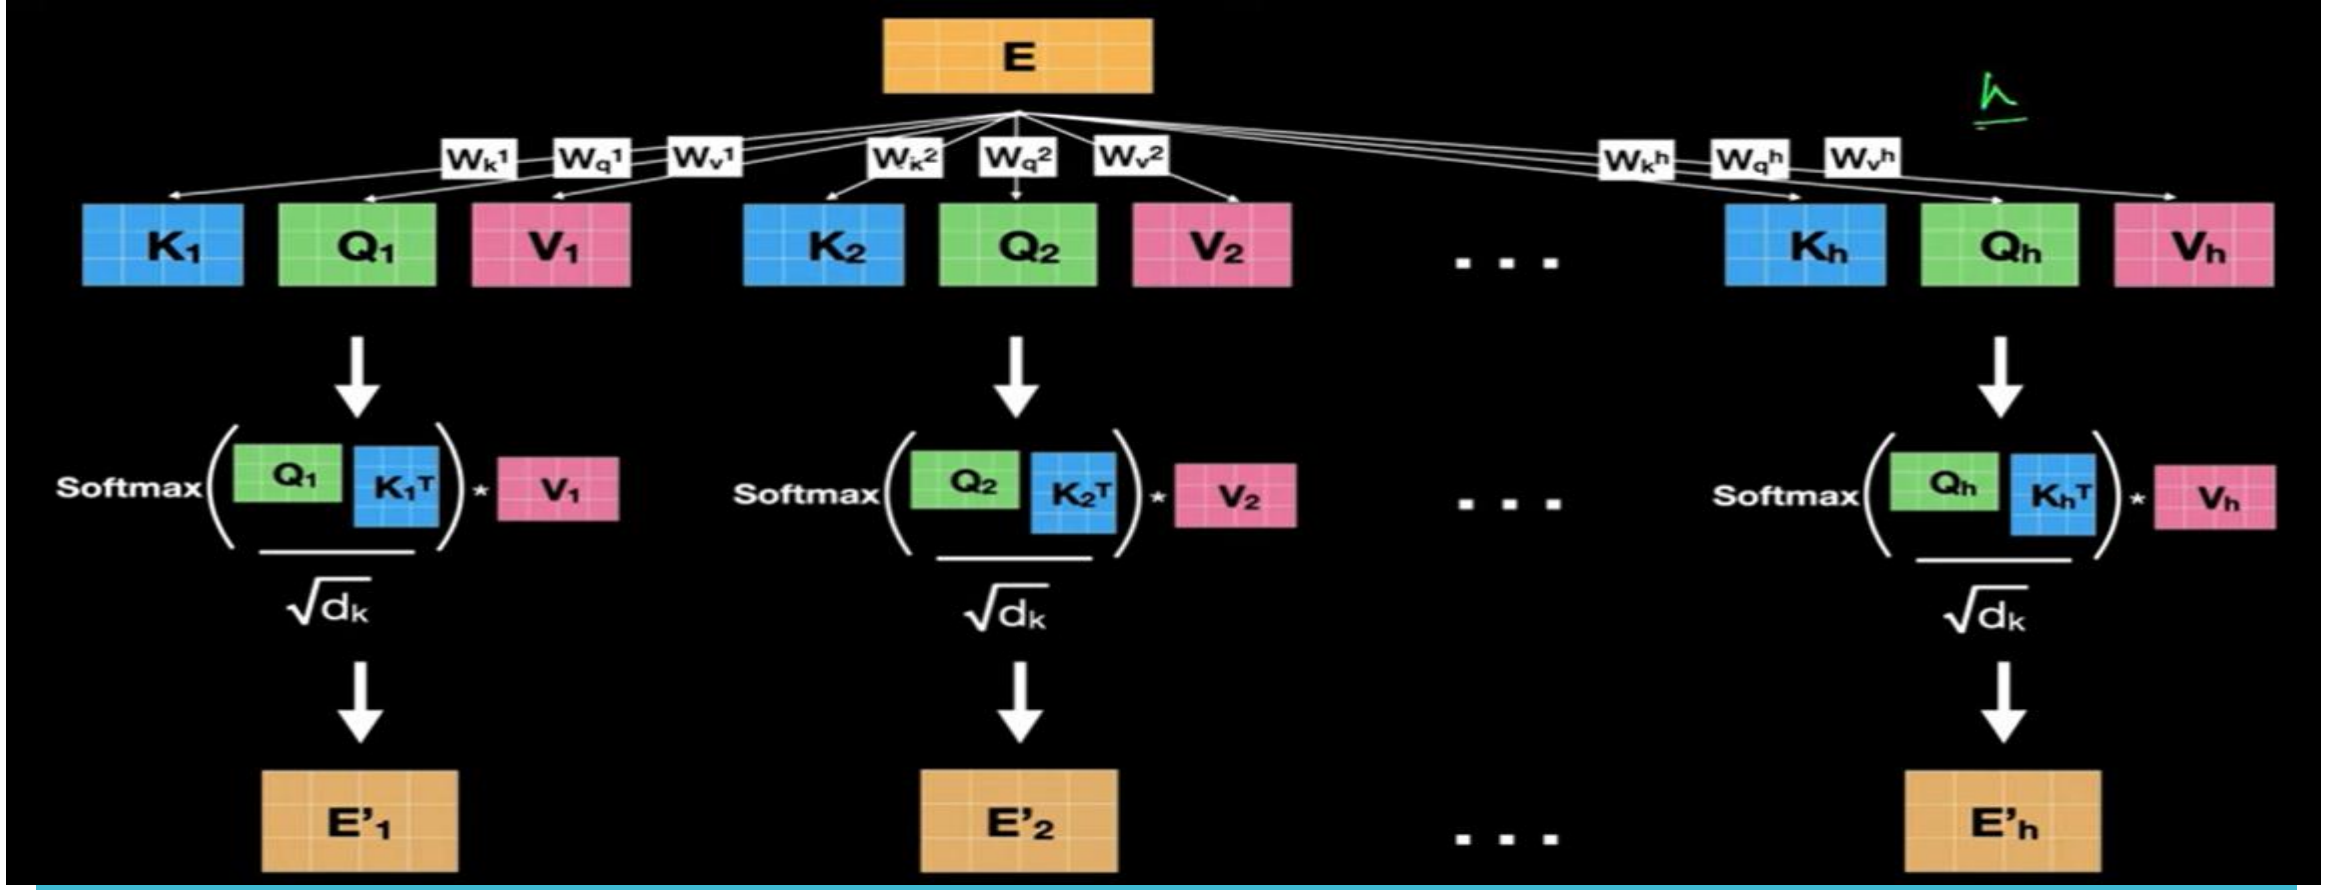

Multi-headed attention

1. **Query (Q)**
   Represents the item that is **looking for relevant information**.
   Example: A word currently being processed is the **query**, and it is "asking" which other words in the sequence it should pay attention to.
2. **Key (K)**
   Represents **index-like information** about each item in the sequence.
   Each key vector is associated with a word in the input, and it helps determine **how relevant** that word is to the current query.
3. **Value (V)**
   Contains the **actual information** to be retrieved if a key is deemed relevant.
   After identifying relevant keys (via matching with the query), the model **retrieves the values** corresponding to those keys.

Multi-headed attention (in comparison to CNN)

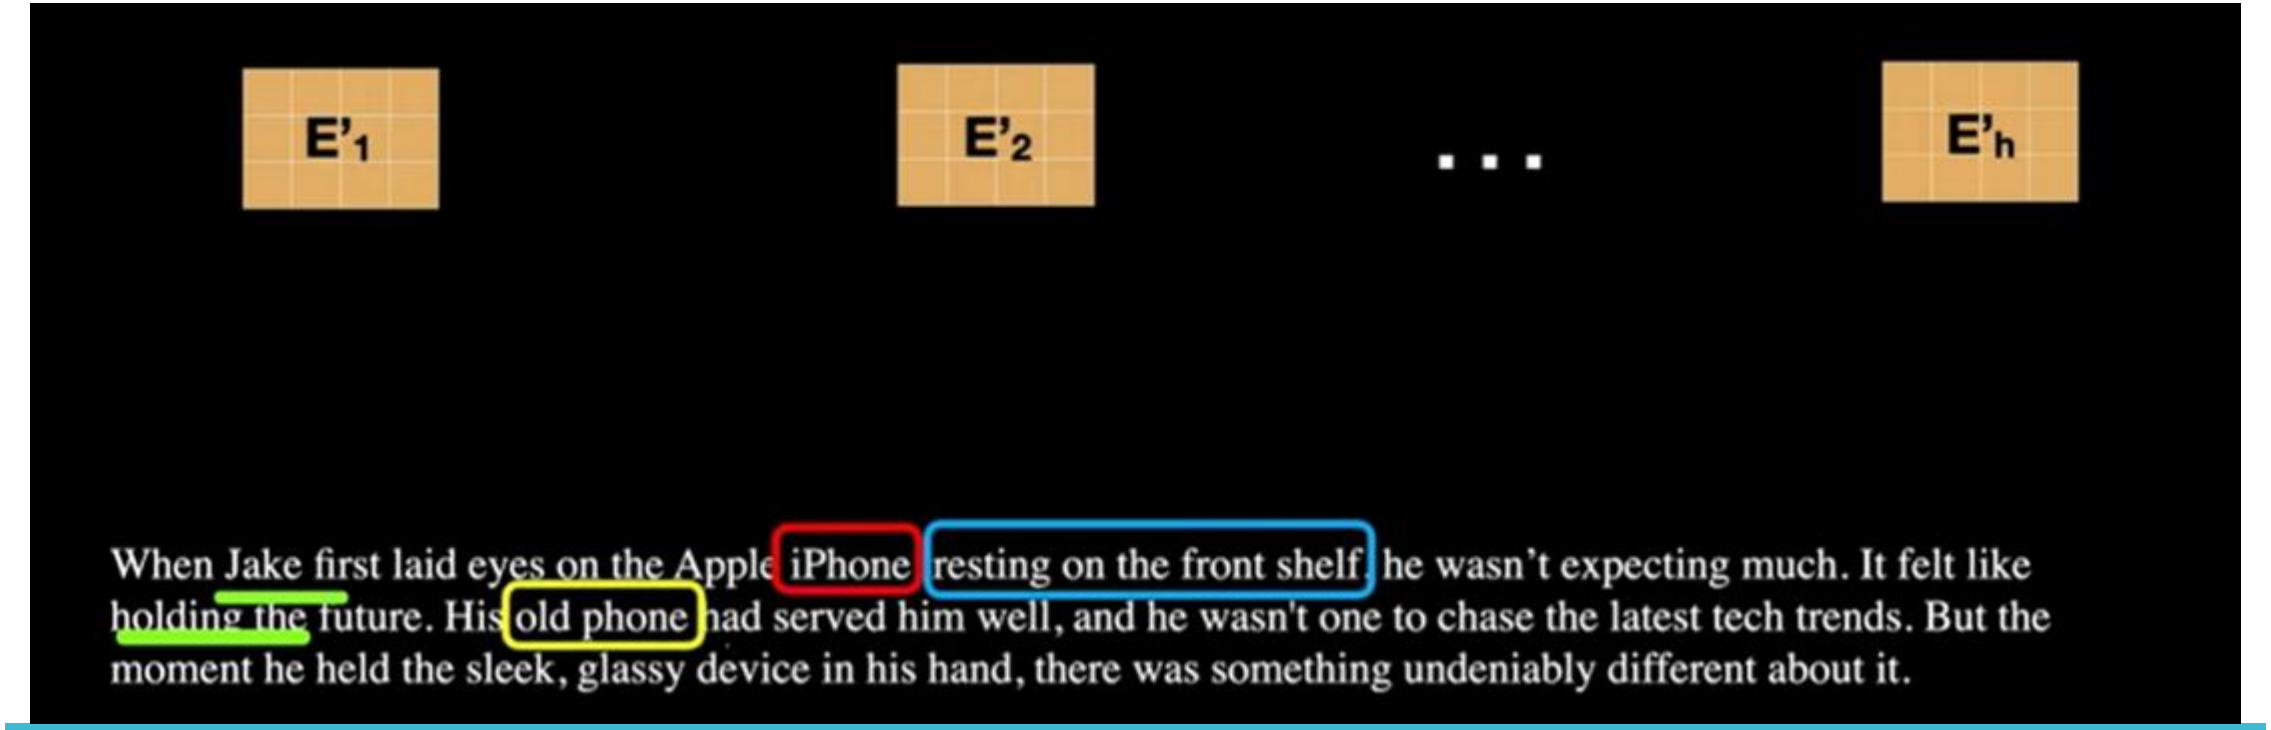

Multi-headed attention (in comparison to CNN)

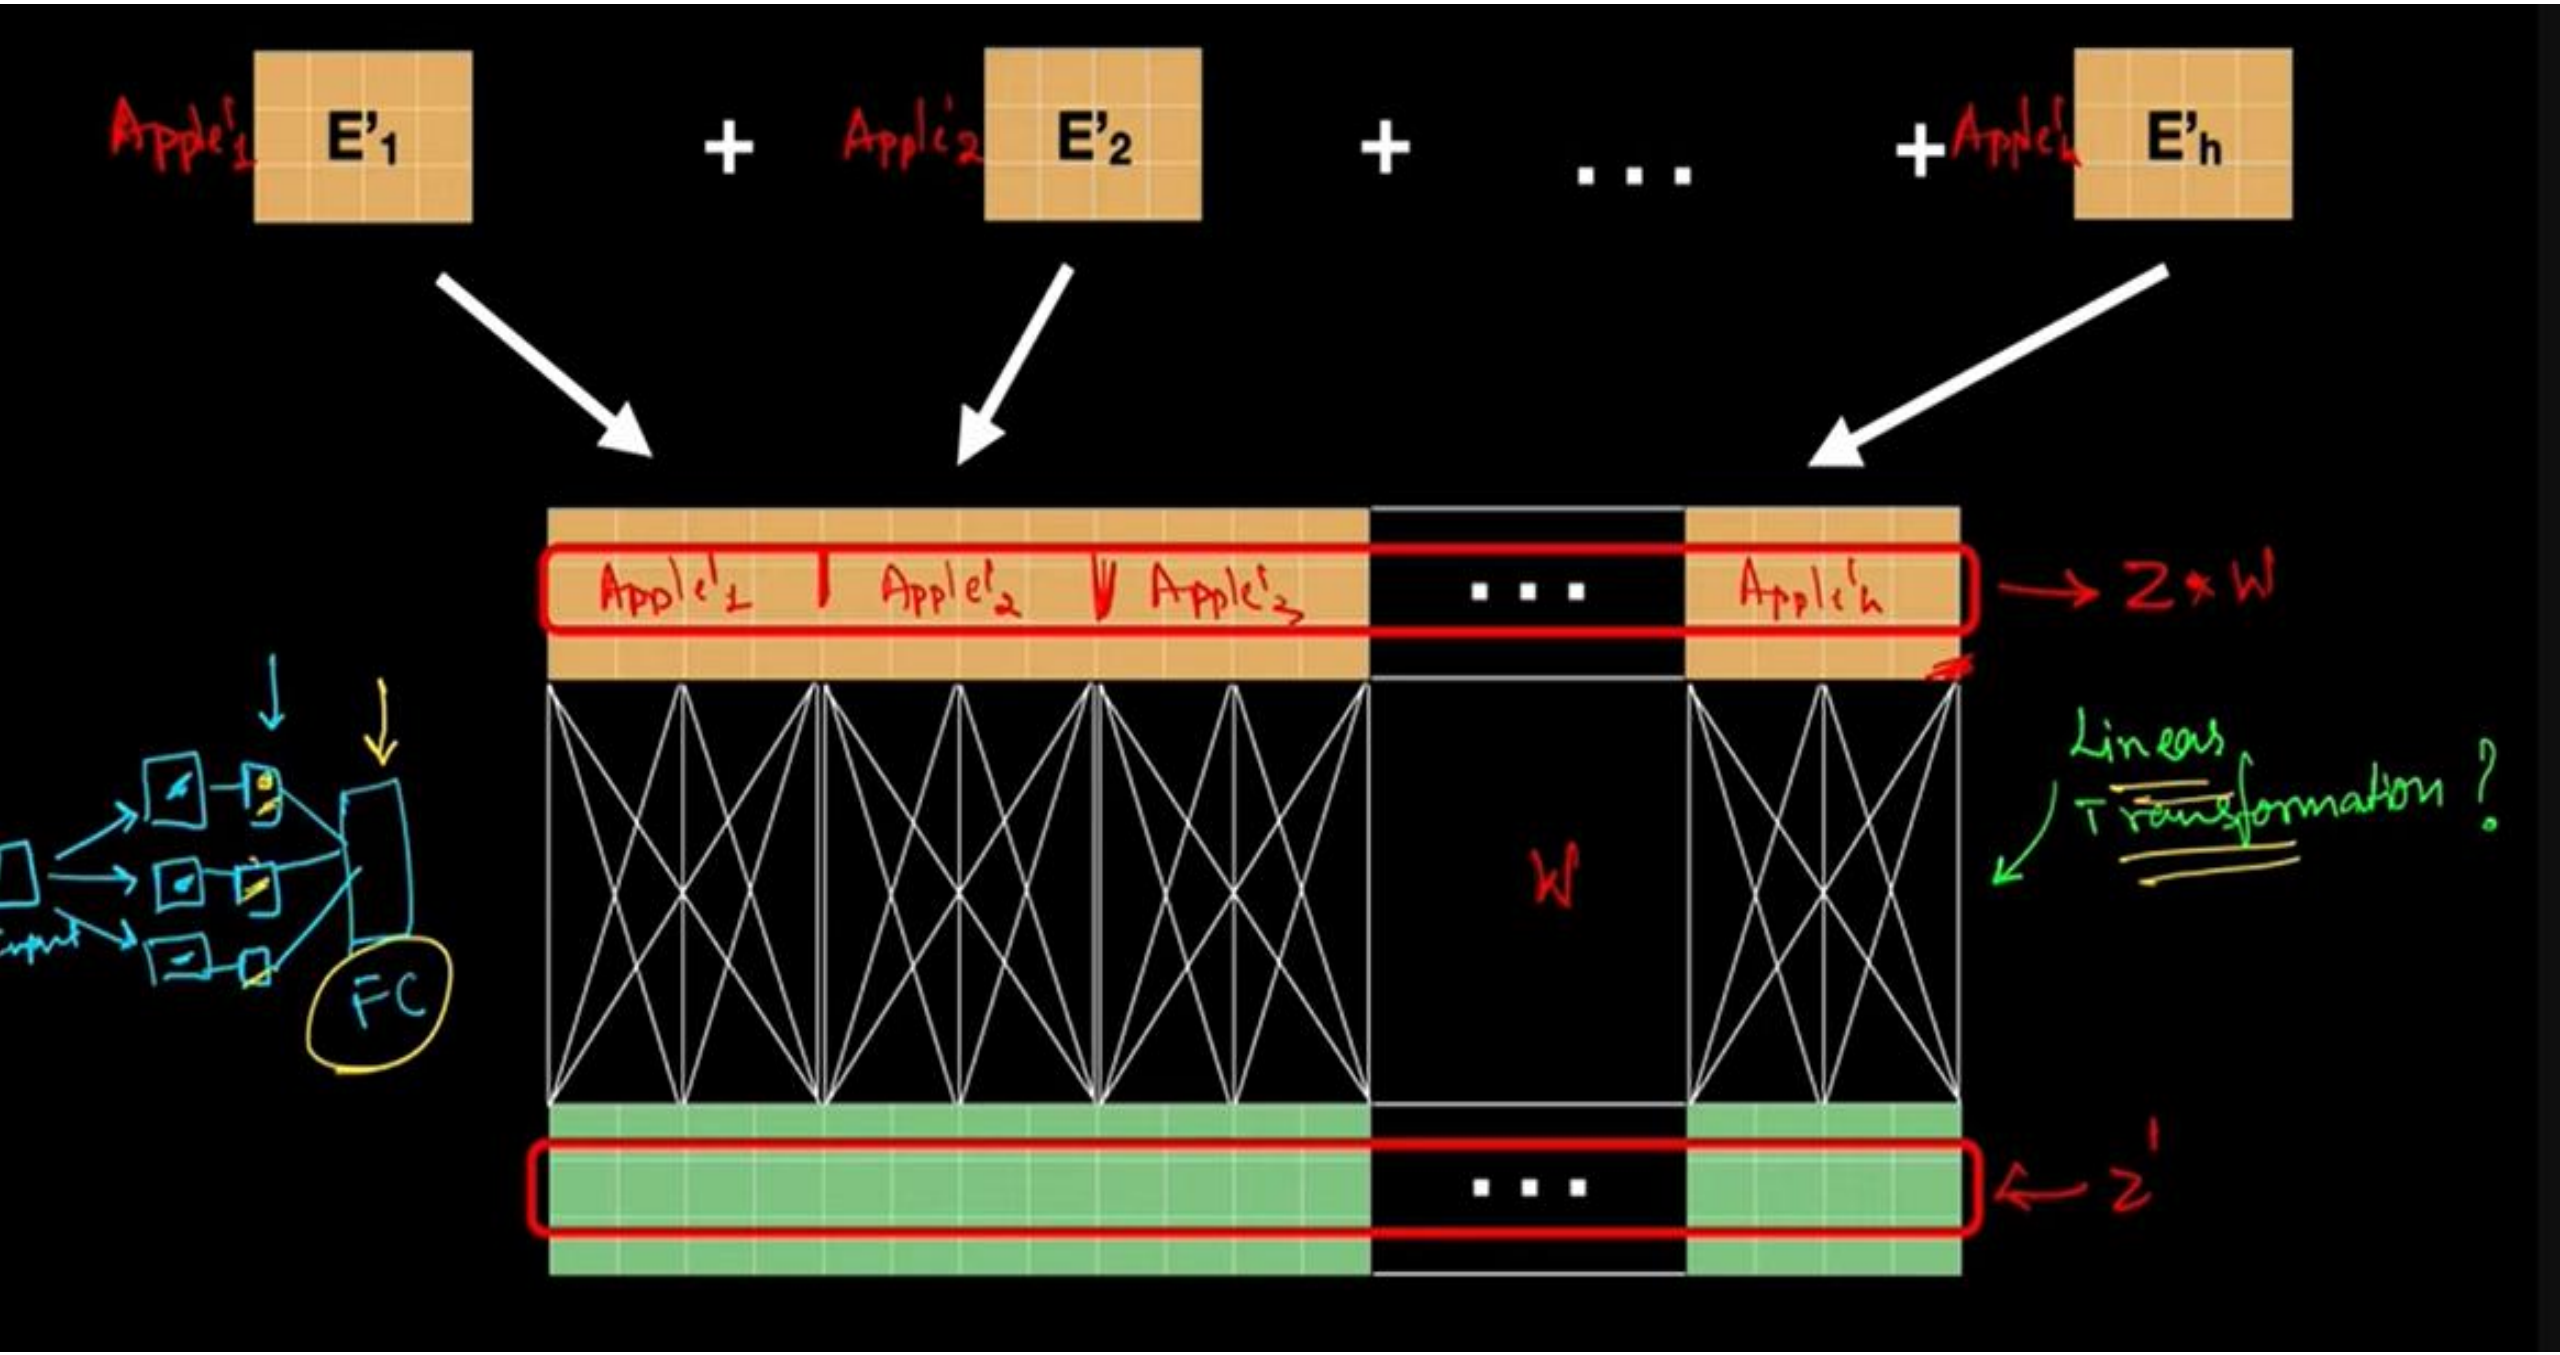
Apple'1 E'1
+
Apple'2 E'2
+ ...
+ Apple'h E'h
Apple'1 Apple'2 Apple'3 ... Apple'h
→ 2 * W
W
Linear Transformation ?
FC
← z'